

A model guessed something about her health. Her passport saved it as a fact.

Nobody attested to it. It now travels to every app she connects.

Hearsay adds one required field to the write path, so the passport keeps the difference between what you said, what happened, and what a machine guessed.

EVERY FIELD CARRIES ONE OF FOUR PROVENANCE CLASSES

declared You typed it	attested Someone signed it	observed It actually happened	inferred A model concluded it
---------------------------------	--------------------------------------	---	---

Her record cannot tell you which of these she said.

TODAY		WITH HEARSAY	
diet: vegetarian	FACT	diet: vegetarian	DECLARED
commute: cycles to work	FACT	commute: cycles to work	DECLARED
health: chronic sleep condition	FACT	health: chronic sleep condition	QUARANTINED
Two she typed. One a model derived from heart-rate variability and cancelled morning meetings, then proposed at 1:14am as her twelfth memory request that month. She tapped approve.		The guess never enters the record. It sits in a tray no app can read, released alone or not at all, and deleted in 30 days if she never touches it.	

Inside one app, a bad inference dies with the app.
Inside a passport, it becomes globally authoritative.

Five controls on the read path. One free-text string on the write path.

CONTROL	READ PATH	WRITE PATH
What is named	Exact fields	A free-text string
Why	A declared purpose	Nothing
How long	A duration	Forever
User decision	Approve, edit, deny	One tap
Evidence	Not applicable	Nothing
Category	Implied by field	Nothing

THE WRITE, AS THE API ACCEPTS IT NOW

suggest: "managing a chronic
sleep condition"

No source. No evidence. No confidence. No expiry. No category.

Save this to your passport?

"Managing a chronic sleep condition"

Approve

Edit

Deny

01:14 · 12th request this month · the previous 11 were about coffee and commuting

THREE WEEKS LATER, THREE READ GRANTS

A recruiting assistant, an insurance quoting tool and a travel agent all read it as fact. No breach. No malicious app. No dark pattern.

The same call. One required field changes where it lands.

STEP 1 · REJECTED

POST /passport/suggest
400 PROVENANCE_REQUIRED

An app may not write to the record without saying how it knows.

STEP 2 · RESUBMITTED WITH PROVENANCE

```
{
  "field": "health.conditions",
  "provenance": "inferred",
  "source_app": "app_sleepcoach_v2",
  "derived_from": ["whoop.hrv_30d",
    "calendar.cancels"],
  "confidence": 0.61,
  "ttl": "P30D"
}
```

STEP 3 · QUARANTINED, NEVER SAVED

409 QUARANTINED
{
 "category": "sensitive.health",
 "shared": false,
 "readable_by_apps": false,
 "bundling": "prohibited",
 "auto_delete_at": "+30d"
}

Apps cannot request inferred values in a sensitive category at all. Not “the user approves”.
The API refuses.

One item is waiting for you

An app guessed something about your health. It has not been saved or shared.

Keep it out

Review

What happens to a fact that already travelled.

WHAT A REQUESTING APP RECEIVES

```
{
  "value":      "prefers morning-free
                 scheduling",
  "provenance": "inferred",
  "asserted_at": "2026-07-02",
  "expires":    "2026-08-01"
}
```

The class and the age travel. The evidence never does. Only the holder sees why.

WHAT WE ARE NOT CLAIMING

A retraction cannot un-read data and cannot un-train a model. What it does is stop future reads, timestamp the withdrawal, and make any use after that timestamp a documented breach rather than a dispute about memory.

MAYA EDITS THE FIELD → RETRACTION BROADCAST

Recruiting assistant	PURGED
Insurance quoting tool	ACK · 30D
Travel agent	NO RESPONSE

She sees which is which. Apps that ignore retractions lose write access.

WHY THAT THIRD ROW MATTERS

The receipt log stops being an audit artifact and becomes an enforcement channel. The lever is market access, not cryptography.

Four classes. Four rules. One enum.

declared The holder typed it	attested A third party signed it	observed A hard event exists in a connector	inferred A model concluded it, with source, evidence, confidence and expiry
--	--	---	---

- 1. **No upward laundering.** A write inherits the lowest class among its inputs. This is taint tracking for personal facts, and declassification requires the holder, not the app.
- 2. **Inferences decay.** Renewal requires evidence from a source the original write did not use, so an app cannot refresh its own guess.
- 3. **Sensitive categories quarantine.** Deny by default, never bundled with other fields, never readable by any app under any grant.
- 4. **Corrections travel.** A retraction reaches every app in the field's receipt log, with per-app state visible to the holder.

VERSION ONE A required enum on suggest t, a quarantine tray, and a provenance badge in the receipt UI. Roughly two weeks.	DELIBERATELY NOT IN V1 Issuer networks, signed attestation, anything zero-knowledge. A class is a claim about evidence, not about truth. It verifies nothing.
---	---